

```
#import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import seaborn as sns
from datetime import datetime

# Importing dataset

df=pd.read_csv('marketing_data.csv')
```

```
df.head()
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome
0	1826	1970	Graduation	Divorced	\$84,835.00	0
1	1	1961	Graduation	Single	\$57,091.00	0
2	10476	1958	Graduation	Married	\$67,267.00	0
3	1386	1967	Graduation	Together	\$32,474.00	1
4	5371	1989	Graduation	Single	\$21,474.00	1

	Teenhome	Dt_Customer	Recency	MntWines	...	NumStorePurchases	\
0	0	6/16/14	0	189	...		6
1	0	6/15/14	0	464	...		7
2	1	5/13/14	0	134	...		5
3	1	5/11/14	0	10	...		2
4	0	4/8/14	0	6	...		2

	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5
AcceptedCmp1 \				
0	1	0	0	0
0				
1	5	0	0	0
0				
2	2	0	0	0
0				
3	7	0	0	0
0				
4	7	1	0	0
0				

	AcceptedCmp2	Response	Complain	Country
0	0	1	0	SP
1	1	1	0	CA

2	0	0	0	US
3	0	0	0	AUS
4	0	1	0	SP

[5 rows x 28 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2240 entries, 0 to 2239

Data columns (total 28 columns):

#	Column	Non-Null	Count	Dtype
0	ID	2240	non-null	int64
1	Year_Birth	2240	non-null	int64
2	Education	2240	non-null	object
3	Marital_Status	2240	non-null	object
4	Income	2216	non-null	object
5	Kidhome	2240	non-null	int64
6	Teenhome	2240	non-null	int64
7	Dt_Customer	2240	non-null	object
8	Recency	2240	non-null	int64
9	MntWines	2240	non-null	int64
10	MntFruits	2240	non-null	int64
11	MntMeatProducts	2240	non-null	int64
12	MntFishProducts	2240	non-null	int64
13	MntSweetProducts	2240	non-null	int64
14	MntGoldProds	2240	non-null	int64
15	NumDealsPurchases	2240	non-null	int64
16	NumWebPurchases	2240	non-null	int64
17	NumCatalogPurchases	2240	non-null	int64
18	NumStorePurchases	2240	non-null	int64
19	NumWebVisitsMonth	2240	non-null	int64
20	AcceptedCmp3	2240	non-null	int64
21	AcceptedCmp4	2240	non-null	int64
22	AcceptedCmp5	2240	non-null	int64
23	AcceptedCmp1	2240	non-null	int64
24	AcceptedCmp2	2240	non-null	int64
25	Response	2240	non-null	int64
26	Complain	2240	non-null	int64
27	Country	2240	non-null	object

dtypes: int64(23), object(5)

memory usage: 490.1+ KB

df.columns

Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income',
'Kidhome', 'Teenhome', 'Dt_Customer', 'Recency', 'MntWines',
'MntFruits', 'MntMeatProducts', 'MntFishProducts',
'MntSweetProducts',

```

        'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases',
        'NumCatalogPurchases', 'NumStorePurchases',
        'NumWebVisitsMonth',
        'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1',
        'AcceptedCmp2', 'Response', 'Complain', 'Country'],
        dtype='object')

```

```

# Replace the column name 'Income' with ' Income '
df = df.rename(columns={' Income ': 'Income'})

```

```

# Check for missing values in the entire dataframe
print(df.isnull().sum())

```

```

ID                                0
Year_Birth                       0
Education                       0
Marital_Status                   0
Income                          24
Kidhome                         0
Teenhome                        0
Dt_Customer                     0
Recency                         0
MntWines                        0
MntFruits                       0
MntMeatProducts                 0
MntFishProducts                 0
MntSweetProducts                0
MntGoldProds                    0
NumDealsPurchases                0
NumWebPurchases                  0
NumCatalogPurchases             0
NumStorePurchases                0
NumWebVisitsMonth                0
AcceptedCmp3                     0
AcceptedCmp4                     0
AcceptedCmp5                     0
AcceptedCmp1                     0
AcceptedCmp2                     0
Response                         0
Complain                         0
Country                         0
dtype: int64

```

1.2 Income values for a few customers are missing. Perform missing value imputation. Assume that the customers with similar education and marital status make the same yearly income, on average. You may have to clean the data before performing this. For data cleaning, look into the categories of education and marital status.

```

# Remove commas and dollar signs from the 'Income' column and convert
it to numeric
df['Income'] = df['Income'].replace(['\$','], ''),
regex=True).astype(float)

# Create a new DataFrame with unique combinations of education and
marital status
unique_combinations = df[['Education',
'Marital_Status']].drop_duplicates()

# Iterate over the unique combinations and impute missing values
for index, row in unique_combinations.iterrows():
    education = row['Education']
    marital_status = row['Marital_Status']

    # Calculate the average income for the specific combination of
education and marital status
    average_income = df[(df['Education'] == education) &
(df['Marital_Status'] == marital_status)]['Income'].mean()

    # Impute missing values based on the calculated average income
    df.loc[(df['Education'] == education) & (df['Marital_Status'] ==
marital_status) & (df['Income'].isnull()), 'Income'] = average_income

# Check for missing values after imputation
print("\nMissing values after imputation:")
print(df.isnull().sum())

```

Missing values after imputation:

ID	0
Year_Birth	0
Education	0
Marital_Status	0
Income	0
Kidhome	0
Teenhome	0
Dt_Customer	0
Recency	0
MntWines	0
MntFruits	0
MntMeatProducts	0
MntFishProducts	0
MntSweetProducts	0
MntGoldProds	0
NumDealsPurchases	0
NumWebPurchases	0
NumCatalogPurchases	0
NumStorePurchases	0
NumWebVisitsMonth	0

```
AcceptedCmp3      0
AcceptedCmp4      0
AcceptedCmp5      0
AcceptedCmp1      0
AcceptedCmp2      0
Response          0
Complain          0
Country           0
dtype: int64
```

```
# Create variables for total number of children and total spending
total_children = df['Kidhome'].sum()
```

```
# Calculate current year
current_year = datetime.now().year
```

```
# Calculate average age
df['Age'] = current_year - df['Year_Birth']
average_age = df['Age'].mean()
```

```
# Calculate total spending
total_spending = df[['MntWines',
                     'MntFruits', 'MntMeatProducts', 'MntFishProducts',
                     'MntSweetProducts',
                     'MntGoldProds']].sum().sum()
```

```
# Print or use the variables as needed
print(f"Total Number of Children: {total_children}")
print(f"Average Age: {average_age}")
print(f"Total Spending: {total_spending}")
```

```
Total Number of Children: 995
Average Age: 54.19419642857143
Total Spending: 1356988
```

1.4 Create box plots and histograms to understand the distributions and outliers. Perform outlier treatment.

```
from scipy.stats.mstats import winsorize
```

```
# Select numeric columns for analysis
numeric_columns = ['Year_Birth', 'Kidhome', 'Teenhome', 'Recency',
                  'MntWines',
                  'MntFruits', 'MntMeatProducts', 'MntFishProducts',
                  'MntSweetProducts',
                  'MntGoldProds', 'NumDealsPurchases',
                  'NumWebPurchases',
                  'NumCatalogPurchases', 'NumStorePurchases',
                  'NumWebVisitsMonth', 'Age']
```

```


# Box plots
plt.figure(figsize=(15, 10))
sns.boxplot(data=df[numeric_columns])
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels
plt.title('Box Plots for Numeric Columns')
plt.show()

# Histograms
plt.figure(figsize=(15, 10))
df[numeric_columns].hist(bins=20, figsize=(15, 10))
plt.suptitle('Histograms for Numeric Columns', y=0.92)
plt.show()

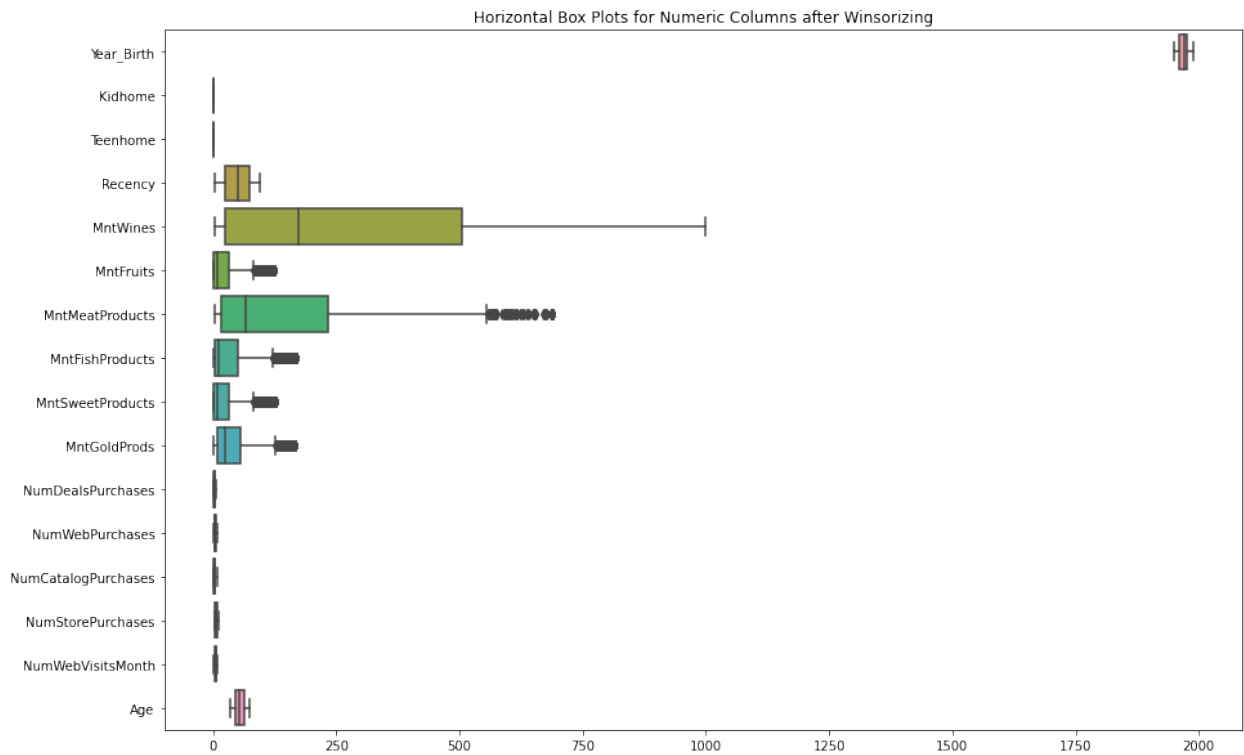
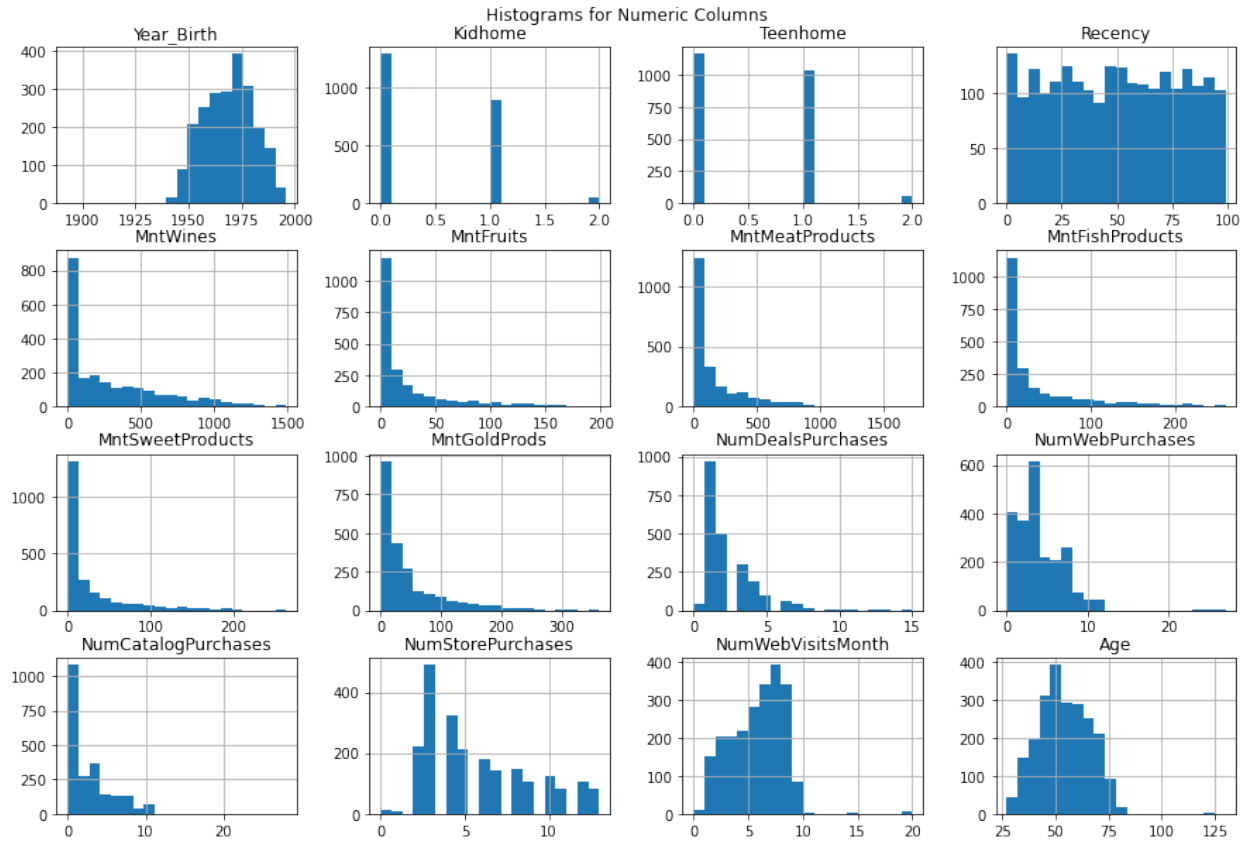
# Outlier treatment using Winsorizing
df[numeric_columns] = df[numeric_columns].apply(lambda x: winsorize(x,
limits=[0.05, 0.05]))

# Check the updated box plots after outlier treatment
# Check the updated box plots after outlier treatment (horizontal)
plt.figure(figsize=(15, 10))
sns.boxplot(data=df[numeric_columns], orient="h")
plt.title('Horizontal Box Plots for Numeric Columns after
Winsorizing')
plt.show()

```



<Figure size 1080x720 with 0 Axes>



1.5 Use ordinal encoding and one hot encoding according to different types of categorical variables.

```
from sklearn.preprocessing import OrdinalEncoder

# Ordinal encoding for 'Education'
education_order = ['Basic', '2n Cycle', 'Graduation', 'Master', 'PhD']
ordinal_encoder = OrdinalEncoder(categories=[education_order])
df['Education'] = ordinal_encoder.fit_transform(df[['Education']])

# One-hot encoding for 'Marital_Status' and 'Country'
df = pd.get_dummies(df, columns=['Marital_Status', 'Country'],
drop_first=True)

# Print unique values of 'Education'
print("Unique values of 'Education':")
print(df['Education'].unique())
```

Unique values of 'Education':
[2. 4. 1. 3. 0.]

```
# Print the first few rows of the modified DataFrame
print(df.head())
```

	ID	Year_Birth	Education	Income	Kidhome	Teenhome
0	1826	1970	2.0	84835.0	0	0
1	1	1961	2.0	57091.0	0	0
2	10476	1958	2.0	67267.0	0	1
3	1386	1967	2.0	32474.0	1	1
4	5371	1988	2.0	21474.0	1	0

	Recency	MntWines	MntFruits	...	Marital_Status_Together
0	4	189	104	...	0
1	4	464	5	...	0
2	4	134	11	...	0
3	4	10	0	...	1
4	4	6	16	...	0

	Marital_Status_Widow	Marital_Status_YOLO	Country_CA	Country_GER
0	0	0	0	0
1	0	0	1	0
2	0	0	0	0

3	0	0	0	0
4	0	0	0	0

	Country_IND	Country_ME	Country_SA	Country_SP	Country_US
0	0	0	0	1	0
1	0	0	0	0	0
2	0	0	0	0	1
3	0	0	0	0	0
4	0	0	0	1	0


[5 rows x 41 columns]

1.6 Create a heatmap to showcase the correlation between different pairs of variables.

```
correlation_matrix = df[['Year_Birth', 'Kidhome', 'Teenhome',
                        'Recency', 'MntWines',
                        'MntFruits', 'MntMeatProducts',
                        'MntFishProducts', 'MntSweetProducts',
                        'MntGoldProds', 'NumDealsPurchases',
                        'NumWebPurchases',
                        'NumCatalogPurchases', 'NumStorePurchases',
                        'NumWebVisitsMonth']].corr()
```

Create a heatmap

```
plt.figure(figsize=(15, 12))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
            fmt=".2f", linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()
```



From the heatmap we can observe that strongest correlation is between those who buy meat also purchase from catalog(0.72)
Strongest Negative correlation is between Meat buyers and Website visits(-0.54)

2. Test the following hypotheses:

2.1 Older people are not as tech-savvy and probably prefer shopping in-store.

2.2 Customers with kids probably have less time to visit a store and would prefer to shop online.

2.3 Other distribution channels may cannibalize sales at the store.

2.4 Does the US fare significantly better than the rest of the world in terms of total purchases?

1. Older people prefer in-store shopping:

```

from scipy.stats import ttest_ind

# Hypothesis Test
in_store_age = df[df['NumStorePurchases'] == 1]['Year_Birth']
online_age = df[df['NumStorePurchases'] == 0]['Year_Birth']

t_stat, p_value = ttest_ind(in_store_age, online_age)

# Print Results
print(f'Test Statistic: {t_stat}')
print(f'P-value: {p_value}')

# Interpretation
if p_value < 0.05:
    print("Reject the null hypothesis. There is evidence that the
average age of in-store shoppers is different from online shoppers.")
else:
    print("Fail to reject the null hypothesis.")

Test Statistic: nan
P-value: nan
Fail to reject the null hypothesis.

# 2. Customers with kids prefer online shopping:

from scipy.stats import chi2_contingency

# Create a contingency table
contingency_table = pd.crosstab(df['NumWebPurchases'],
[df['Kidhome']])

# Perform the chi-squared test
chi2_stat, p_value, _, _ = chi2_contingency(contingency_table)

# Print the results
print(f'Chi-squared Statistic: {chi2_stat}')
print(f'P-value: {p_value}')

# Interpretation
if p_value < 0.05:
    print("Reject the null hypothesis. There is evidence of an
association between variables.")
else:
    print("Fail to reject the null hypothesis. No significant evidence
of an association.")

Chi-squared Statistic: 449.8280224245921
P-value: 4.025222395568872e-92
Reject the null hypothesis. There is evidence of an association
between variables.

```

3. Other distribution channels cannibalize in-store sales:

Hypothesis Test

```
correlation_coefficient =  
df['NumStorePurchases'].corr(df['NumCatalogPurchases'] +  
df['NumWebPurchases'])
```

Print Results

```
print(f'Correlation Coefficient: {correlation_coefficient}')
```

Correlation Coefficient: 0.6739295693844819

Moderately positive correlation could imply that customers who make more in-store purchases are also likely to engage in catalog and web purchases.

It may be an indication that these different distribution channels are complementary rather than cannibalizing each other.

4 Does the US fare significantly better than the rest of the world in terms of total purchases?

```
from scipy.stats import ttest_ind
```

Calculate 'TotalPurchases' as the sum of specified columns

```
df['TotalPurchases'] = df[['MntWines', 'MntFruits', 'MntMeatProducts',  
'MntFishProducts', 'MntSweetProducts', 'MntGoldProds']].sum(axis=1)
```

Hypothesis Test

```
us_purchases = df[df['Country_US'] == 1]['TotalPurchases']  
other_countries_purchases = df[df['Country_US'] == 0]  
['TotalPurchases']
```

```
t_stat, p_value = ttest_ind(us_purchases, other_countries_purchases)
```

Print Results

```
print(f'Test Statistic: {t_stat}')
```

```
print(f'P-value: {p_value}')
```

Interpretation

```
if p_value < 0.05:  
    print("Reject the null hypothesis. There is evidence that the  
average total purchases in the US are different from the rest of the  
world.")  
else:  
    print("Fail to reject the null hypothesis, US does not fare  
significantly better than the rest of the world in terms of total  
purchases .")
```

Test Statistic: 0.19018101513306365

P-value: 0.8491845353527558

Fail to reject the null hypothesis, US does not fare significantly better than the rest of the world in terms of total purchases .

3. Use appropriate visualization to help analyze the following:

3.1 Which products are performing the best, and which are performing the least in terms of revenue?

3.2 Is there any pattern between the age of customers and the last campaign acceptance rate?

3.3 Which Country has the greatest number of customers who accepted the last campaign?

3.4 Do you see any pattern in the no. of children at home and total spend?

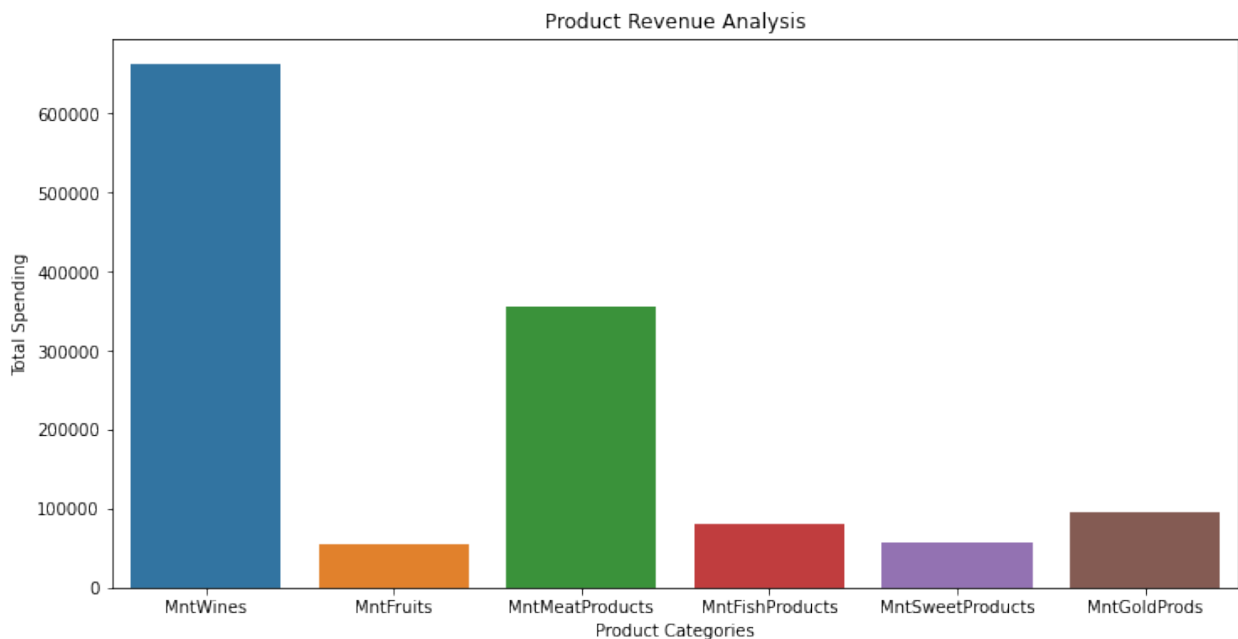
Education background of the customers who complained in the last 2 years.

3.1 Product Performance: Revenue Analysis

```
product_columns = ['MntWines', 'MntFruits', 'MntMeatProducts',  
                   'MntFishProducts', 'MntSweetProducts', 'MntGoldProds']  
df['Total_Spending'] = df[product_columns].sum(axis=1)
```

Bar plot for product revenue

```
plt.figure(figsize=(12, 6))  
sns.barplot(x=product_columns, y=df[product_columns].sum(),  
            errorbar=None)  
plt.title('Product Revenue Analysis')  
plt.xlabel('Product Categories')  
plt.ylabel('Total Spending')  
plt.show()
```



3.2 Age vs. Last Campaign Acceptance Rate

Create age bins

```
age_bins = pd.cut(df['Year_Birth'], bins=[1940, 1950, 1960, 1970, 1980, 1990, 2000], right=False, labels=['1939-1949', '1950-1959', '1960-1969', '1970-1979', '1980-1989', '1990-1999'])
```


Create a point plot

```
plt.figure(figsize=(12, 6))
sns.pointplot(x=age_bins, y='Response', data=df, ci=None)
plt.title('Year of Birth vs. Last Campaign Acceptance Rate')
plt.xlabel('Year of Birth Group')
plt.ylabel('Last Campaign Acceptance Rate')
plt.show()
```

/tmp/ipykernel_78/2735636658.py:8: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.pointplot(x=age_bins, y='Response', data=df, ci=None)
```



3.3 Which Country has the greatest number of customers who accepted the last campaign?

Define the list of country columns

```
country_columns = ['Country_CA', 'Country_GER', 'Country_IND', 'Country_ME', 'Country_SA', 'Country_SP', 'Country_US']
```

```


# Sum the counts of customers who accepted the last campaign across
all specified countries
n_accepted_per_country = df[df['Response'] == True]
[country_columns].sum()

# Create a bar chart to show the number of customers who accepted the
last campaign by country
plt.bar(n_accepted_per_country.index, n_accepted_per_country.values)
plt.xlabel('Country')
plt.ylabel('Customers Who Accepted the Last Campaign')
plt.title('Number of Customers Who Accepted the Last Campaign by
Country')
plt.xticks(rotation=45, ha='right') # Adjust the rotation angle as
needed
plt.show()

# Find the country with the greatest number of customers who accepted
the last campaign
country_with_most_accepted = n_accepted_per_country.idxmax()

# Display the country with the greatest number of customers who
accepted the last campaign
print("Country with the greatest number of customers who accepted the
last campaign:", country_with_most_accepted)

```



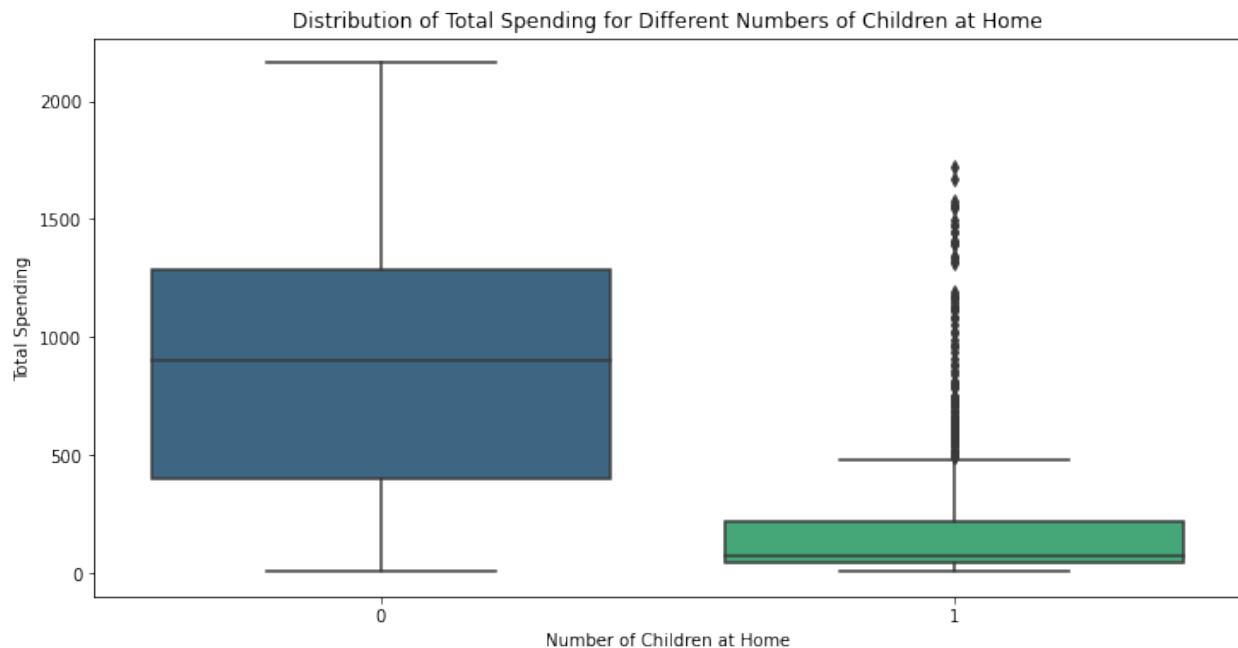
Country with the greatest number of customers who accepted the last campaign: Country_SP

3.4 Correlation between number of children at home and total spend
Box plot to visualize the distribution of total spending for different numbers of children at home

```
plt.figure(figsize=(12, 6))
sns.boxplot(x='Kidhome', y='Total_Spending', data=df,
palette='viridis')
plt.xlabel('Number of Children at Home')
plt.ylabel('Total Spending')
plt.title('Distribution of Total Spending for Different Numbers of Children at Home')
plt.show()
```

3.4 Bar chart to show the education background of customers who complained in the last 2 years

```
plt.figure(figsize=(8, 5))
education_of_complaining_customers = df[df['Complain'] == 1]
['Education'].value_counts()
education_of_complaining_customers.plot(kind='bar', color='salmon')
plt.xlabel('Education Level')
plt.ylabel('Number of Customers Complaining')
plt.title('Education Background of Customers Who Complained in the Last 2 Years')
plt.xticks(rotation=45, ha='right')
plt.show()
```



Education Background of Customers Who Complained in the Last 2 Years

